


How much of a DSA attention block is the top-k search? H100 PCIe, k=2048, eager
block inputs are prepared projections — no q/k projections, RoPE, output projection or MLP

GPU latency (μ s) — bars sum separately measured stages

DSA attention block, stage composition



Search vs attention, as a share of the block

