


MLA and FFN, and where each top-k search sits inside them — dense layer (3 of 61) H100 PCIe, k=2048, cudagraph
 DSA is the selection front-end inside MLA, not MLA itself · one layer of 61 · experts as per-(token,expert) GEMVs, routing fixed

The layer as two sub-layers
 black rule = MLA/FFN boundary · top number = layer total, ms



MLA grows with context, FFN does not
 the only context-dependent work in the layer is the DSA search

