

Every GEMM in the layer runs near the memory roofline; the selection kernels run 200-10000x below it H100 PCIe, cudagraph bytes are compulsory traffic, so these are upper bounds on efficiency. Small elementwise stages (post-norm) are latency-bound too, but cost little time — selection is both inefficient and expensive.

