

The same layer with expert execution excluded — dense layer (3 of 61) H100 PCIe, cudagraph
NOT the full layer: the feed-forward removed here is 41–48% of it (see Figure 9 for the complete picture)

