

The same layer with expert execution excluded — MoE layer (58 of 61) H100 PCIe, eager
 NOT the full layer: the feed-forward removed here is 41-88% of it (see Figure 9 for the complete picture)

