


Profile of the MLA sub-layer, block by block — MoE layer, H100 PCIe, k=2048, cudagraph
FFN excluded · DSA is inside MLA · the output projection is the single largest block, attention the smallest

Every block of the MLA sub-layer, N=128K



Rolled into families, across context lengths
x ticks are cached tokens N · number above bar = MLA total, μs

